

# Uninformed Search: Problem-solving Agents & Path Planning

---

CS3243: Introduction to Artificial Intelligence – Lecture 2

17 January 2022

# Contents

---

1. Administrative Matters
2. Problem-Solving Agents
3. Search Spaces
4. Search Solutions
5. Breadth-First Search
5. Uniform-Cost Search
6. Depth-First Search
7. Depth-Limited and Iterative Deepening Search
8. Tree-Search Versus Graph-Search

# Administrative Matters

---

# Upcoming...

---

- Tutorials

- Increased to 16 classes
- Begin next week (i.e., Week 3)
- Face-to-face in SR-2

CS3243

296 students

- Deadlines

- DQ0, DQ1 (released last Monday)
  - *Due this Friday (21 January), 2359 hrs*
- DQ2 (released today)
  - *Due this Friday (21 January), 2359 hrs*
- TA1 (released today)
  - *Due this Sunday (23 January), 2359 hrs*
  - *Refer to the tutorial assignment instructions document on LumiNUS*

# Project Poll

## Grade distribution

- 20% competitive - i.e., driven by performance of the solutions submitted by your peers
  - Specifically,  $\text{percentage\_obtained} = 20 * [(\text{total\_students} - (\text{your\_rank} - 1)) / \text{total\_students}]$
  - Ranks based on 1224 ranking system
- 30% tough cases - i.e., some optimisations would be required to clear
- 50% applications on standard difficulty - i.e., basic implementations

Note that only the 20% component is different between the two options.

24%

Version 1:  
20% on competitive component

24%



156 responses

## Grade distribution

- 20% hidden test cases - i.e., hidden cases requiring several optimisations
- 30% tough cases - i.e., some optimisations would be required to clear
- 50% applications on standard difficulty - i.e., basic implementations

Note that only the 20% component is different between the two options.

76%

Version 2:  
20% on hidden test cases

76%

# Administrative Questions Before We Begin?

---

- Any pressing questions about the administrative matters?
- Channels
  - Verbally on Zoom
  - On Archipelago
  - Via Zoom Chat



OR <https://archipelago.rocks/app/resend-invite/54568232609>

# Problem-Solving Agents

---

# Recall from Lecture 1...

- Goal in AI → determine agent function  $f$

- $f: P \rightarrow a$

- $a \in A$

- Key idea → AI as graph search

- Each percept corresponds to a state in the problem (state → vertex)

- Define the desired states → goals

- After each action, we arrive at a new state (action → edge)

- Construct a search space (graph)

- Design and apply a graph search algorithm

(1) Define performance measure and search space

(2) Design search algorithm


Goal-based Agent



# Problem-Solving Agent: Example Application

- Consider a Maze Puzzle problem

- Layout known
- Moves  $\leftarrow, \uparrow, \rightarrow, \downarrow$
- Find path from **S** to **G**



- Graph formulation


- Vertices (states): positions in maze
- Edges (actions): moves (e.g., b0 to b1 – i.e.,  $\downarrow$ )

Summary of Adjacency Matrix

State	Actions
b0 (S)	b1
b1	b0, b2, c1
b2	b1, b3
b3	b2, c3
c1	b1, d1
c3	b3
d1	c1, e1
e1	d1, e2
e2	e1, e3
e3	e2, f3
f3 (G)	e3

# Problem-Solving Agent: Example Application

- Resultant graph



- Solve using graph search algorithm

Summary of Adjacency Matrix

State	Actions
b0 (S)	b1
b1	b0, b2, c1
b2	b1, b3
b3	b2, c3
c1	b1, d1
c3	b3
d1	c1, e1
e1	d1, e2
e2	e1, e3
e3	e2, f3
f3 (G)	e3

# Path Planning Problem Properties

---

- Assumed environment

- Fully observable
- Deterministic
- Discrete
- Episodic

## *Episodic interpretation?*

- Static + complete information
- Taking actions only changes your position
- Able to **PLAN** → look ahead at what to do
- Execute plan once defined

- Plan is formed sequentially

- Each action in the plan impacts the next action in the plan
- Development of the one plan has no bearing on the next → episodic problem

# Search Spaces

---

# Search Space Definition

---

- State representation,  $s_i$ 
  - ADT containing data describing an instance of the environment
  - Initial state ( $s_0$ )
- Goal test, *isGoal*:  $s_i \rightarrow \{0, 1\}$ 
  - Function that returns 1 if given state  $s_i$  is a goal state, else returns 0
- Actions, *actions*:  $s_i \rightarrow A$ 
  - Function that returns the possible actions,  $A = \{a_1, \dots, a_k\}$ , at a given state,  $s_i$
- Action costs, *cost*:  $(s_i, a_j, s_i') \rightarrow v$ 
  - Function that returns cost,  $v$ , of taking the action  $a_j$  at state  $s_i$  to reach state  $s_i'$
  - Generally, assume costs  $\geq 0$

# Search Space Definition

- Transition model,  $T: (s_i, a_j) \rightarrow s_i'$ 
  - Function that returns the state transitioned to,  $s_i'$ , when action  $a_j$  is applied at state  $s_i$
- ***Generally applicable*** to many AI problems (with slight modifications)
- Actions / transition model / action costs functions
  - Potential ***search space generalisation***
  - Example
    - Transition model:
      - $\leftarrow = (r, c-1)$
      - $\uparrow = (r-1, c)$
      - $\rightarrow = (r, c+1)$
      - $\downarrow = (r+1, c)$
    - Obstacle hash table,  $O$  (assuming less obstacles than valid adjacency cells)
    - Map dimensions
    - Actions function determines non-blocked moves
    - Consider  $10^3$  by  $10^3$  grid with no obstacles
      - Adj. matrix size  $10^6$
      - $|O| = 0$

# Search Solutions

---

# General Search Algorithm

---

```
frontier = {initial state} // frontier is a data structure
while frontier not empty:
    current = frontier.pop()
    if isGoal(current) return path found
    for a in actions(current):
        frontier.push(T(current, a))
return failure
```

- Frontier
  - Part of the search space we are exploring
  - Current edge of the search tree
- Note that each element of the frontier must include
  - Referenced state  $s$
  - Path that was taken to get to  $s$



# States Versus Nodes

---

- State
    - A representation of the environment at some timestamp
  - Node
    - Element in the frontier representing current path traversed
    - Includes the following information
      - State
      - Parent node – to track current path from initial state
      - Action
      - Path cost
      - Depth
- } we will see why we need these later

# Uninformed Search Algorithms

---

- Uninformed → no domain knowledge beyond search problem formulation
- Algorithm differences largely based on frontier implementation
  - Breadth-First Search (BFS): frontier = queue
  - Uniform-Cost Search (UCS): frontier = priority queue
  - Depth-First Search (DFS): frontier = stack
  - Depth-Limited Search (DLS): variation of DFS with max depth
  - Iterative Deepening Search (IDS): iterative version of DLS

# Algorithm Criteria

---

- Efficiency
  - Time complexity
  - Space complexity
- Correctness
  - An algorithm is ***complete*** if it will find a solution when one exists and correctly report failure when it does not
  - An algorithm is ***optimal*** if it finds a solution with the lowest path cost among all solutions (i.e., path cost optimal)

# Questions on the Lecture so far?

---

- Was anything unclear?
- Do you need to clarify anything?
- Channels
  - Verbally on Zoom
  - On Archipelago
  - Via Zoom Chat



OR <https://archipelago.rocks/app/resend-invite/54568232609>

# Breadth-First Search

---

# Breadth-First Search (BFS) Algorithm: A Recap

- Frontier: Queue
- Time Complexity:  $O(b^d)$
- Space Complexity:  $O(b^d)$
- Complete: Yes<sup>1</sup>
- Optimal: No<sup>2</sup>


b: branching factor

d: depth of shallowest goal

1: if (i) b finite AND (ii) state space finite or contains solution

2: optimal if costs uniform (and some other cases)

Example:



Tie-breaking: alphabetic  
order on push to frontier

Frontier Trace:

ITR1 = [A(-)]

ITR2 = [B(A), C(A)]

ITR3 = [C(A), D(A,B), E(A,B)]


ITR4 = [D(A,B), E(A,B), F(A,C), G(A,C)]

ITR5 = [E(A,B), F(A,C), G(A,C)]

ITR6 = DONE (A,B,E)


# BFS: A Simple Improvement

- Assume traversal from left to right
- Will add all denoted nodes to frontier before checking  $G$
- Nodes on level  $d \geq$  sum of all nodes in previous levels assuming  $b \geq 2$



- Performing goal test on pushing to frontier instead of popping from frontier will prevent this with no change in the BFS solution

*Early Goal Test* (as opposed to the original *Late Goal Test*)



# Uniform-Cost Search

---



# Uniform-Cost Search (UCS) Algorithm: A Recap

UCS is basically Dijkstra's Algorithm

- Frontier: Priority Queue<sup>1</sup>
- Time Complexity:  $O(b^e)$
- Space Complexity:  $O(b^e)$
- Complete: Yes<sup>2</sup>
- Optimal: Yes


1: prioritising lower path cost,  $g(n)$ , where  $g(n)$  = path cost of the path taken to reach  $n$

$b$ : branching factor

$e$ :  $1 + \lfloor C^* / \epsilon \rfloor$ , where  $C^*$  is the optimal path cost and  $\epsilon$  is some small positive constant

2: requires same completeness criteria as BFS and that actions costs are  $> \epsilon > 0$

Example:



Tie-breaking: nodes ordered alphabetically when priority is the same

Frontier Trace:

ITR1 = [A((-),0)]

ITR2 = [C((A),1), B((A),3)]

ITR3 = [B((A),3), F((A,C),3), G((A,C),6)]

ITR4 = [F((A,C),3), E((A,B),4), D((A,B),6), G((A,C),6)]


ITR5 = DONE (A,C,F)

Note:

Updating path cost from each node is  $O(1)$  since we store the current path cost


# Why $O(b^e)$ Complexity for UCS?

- UCS explores all paths radiating from the initial state
- UCS explores paths in increments of  $\epsilon$  (smallest action cost)
  - With each *step* it extends paths by at least  $\epsilon$  (from the initial state)
  - Considers paths with cost 0 (initial state only), then cost  $\epsilon$ , then  $2\epsilon$ , etc.
  - Expected to reach goal in  $\lfloor C^* / \epsilon \rfloor + 1$  steps, where  $C^*$  is the optimal path cost



# Why is UCS Optimal? A General Idea

- UCS traverses paths in order of path cost
  - This is because path costs from the initial state are always increasing (given  $\epsilon$ )
    - i.e., whenever a node,  $n$ , is added to a path,  $P$ , the new path,  $P'$  must have a path cost that is at least  $\epsilon$  greater than the past cost of  $P$
- UCS finds the optimal path to each node
  - Suppose UCS outputs path  $P = Q + P'$  as the solution for  $s$  to  $t$
  - Suppose the optimal path from  $s$  to  $t$  is instead  $T = Q + T'$



- UCS must skip shorter paths between  $k$  and  $t$  for it to have chosen  $P$ , which is a contradiction since it always chooses shorter paths to explore first

# Depth-First Search

---

# Depth-First Search (DFS) Algorithm: A Recap


- Frontier: Stack
- Time Complexity:  $O(b^m)$
- Space Complexity:  $O(bm)$
- Complete: No
- Optimal: No

b: branching factor  
m: maximum depth

*Why is DFS incomplete (even under the same assumptions of completeness for BFS)?*

*DFS might get caught in a cycle*

Example:



Tie-breaking: reverse alphabetic order on push to frontier

Frontier Trace:

ITR1 = [A(-)]

ITR2 = [B(A), C(A)]

ITR3 = [D(A,B), E(A,B), C(A)]

ITR4 = [E(A,B), C(A)]

ITR5 = DONE (A,B,E)

Note:

Space efficiency may be improved to  $O(m)$  by simply backtracking – i.e., tracing back to parent and last action taken (assuming fixed order of actions – recall that we store parent node and action taken at parent)

# Depth-Limited & Iterative Deepening Search

---

# Depth-Limited Search (DLS)

---

- DFS with a depth limit,  $\ell$ 
  - Search only up to depth  $\ell$
  - Assume no actions may be taken from nodes at depth  $\ell$
- Same guarantees as DFS with  $\ell$  in place of  $m$ 
  - Time complexity:  $O(b^\ell)$
  - Space complexity:  $O(b\ell)$
  - Complete: No
  - Optimal: No

# Iterative Deepening Search

---

- Idea: use DLS iteratively, each time increasing  $\ell$  by 1
  - Will be completely search based on depth
  - Completeness of BFS with space complexity of DFS
- Overheads: will rerun top levels many times
  - Assuming branching factor  $b$  and depth  $\ell$ , nodes generated by DLS:
    - $O(b^0) + O(b^1) + O(b^2) + \dots + O(b^{\ell-2}) + O(b^{\ell-1}) + O(b^\ell)$
  - Nodes generated by IDS to depth  $d$  with branching factor  $b$ :
    - $(d + 1)O(b^0) + dO(b^1) + (d - 1)O(b^2) + \dots + 3O(b^{d-2}) + 2O(b^{d-1}) + O(b^d)$



# Iterative Deepening Search

---

- Example,  $b = 10$  and  $d = 5$ 
  - $N_{DLS} = 1 + 10 + 100 + 1,000 + 10,000 + 100,000 = 111,111$
  - $N_{IDS} = 6 + 50 + 400 + 3,000 + 20,000 + 100,000 = 123,456$
  - Overhead  $\approx 11\%$
- IDS properties
  - Time:  $O(b^d)$
  - Space:  $O(bd)$
  - Complete: Yes ( $b$  finite and  $d$  finite or contains solution) – same as BFS
  - Optimal: No (optimal if costs uniform (and some other cases)) – same as BFS

# Summary

## ■ Performance under tree-search

Criterion	BFS	UCS	DFS	DLS	IDS
Complete?	Yes <sup>1</sup>	Yes <sup>1,2</sup>	No	No	Yes <sup>1</sup>
Optimal Cost?	Yes <sup>3</sup>	Yes	No	No	Yes <sup>3</sup>
Time	$O(b^d)$	$O(b^{1 + \lceil C^* / \epsilon \rceil})$	$O(b^m)$	$O(b^\ell)$	$O(b^d)$
Space	$O(b^d)$	$O(b^{1 + \lceil C^* / \epsilon \rceil})$	$O(bm)$	$O(b\ell)$	$O(bd)$

1. Complete if  $b$  finite and state space either finite or has a solution

2. Complete if all actions costs are  $> \epsilon > 0$

3. Cost optimal if action costs are all identical (and several other cases)

- Recall that an Early Goal Test on BFS may improve runtime practically
- UCS must perform a Late Goal Test to be optimal (this also accounts for the +1 in the index of its complexity)
- DFS is not complete (even under 1) as it might get caught in a cycle
- DFS space complexity may be improved to  $O(m)$  with backtracking (similar for DLS and IDS)

# Tree-Search Versus Graph-Search

---

# Cycles & Redundant Paths

---

- Cycle → cyclic graph
  - Infinite loops (incomplete)
  - May greatly increase necessary computation
- Redundant path to  $s_i$  → more expensive paths from  $s_0$  to  $s_i$ 
  - Should not use these in solution if optimality is required
- Typical practice → graph-search implementation
  - Maintain a *reached* hash table
  - Add nodes corresponding to each state reached
  - Only add new node to *frontier* (and *reached*) if
    - state represented by node not previously reached
    - path to state already reached is cheaper than one stored

Alternative is tree-search implementation → allow revisits (all we reviewed earlier was done under tree-search)

# Graph-Search Implementations

- Performance under graph-search

Criterion	BFS	UCS	DFS	DLS	IDS
Complete?	Yes <sup>1</sup>	Yes <sup>1,2</sup>	Yes <sup>1</sup>	No	Yes <sup>1</sup>
Optimal Cost?	Yes <sup>3</sup>	Yes	No	No	Yes <sup>3</sup>
Time	$O( V  +  E )$				
Space					

1. Complete if  $b$  finite and state space either finite or has a solution

2. Complete if all actions costs are  $> \epsilon > 0$

3. Cost optimal if action costs are all identical (and several other cases)

- DFS under graph search is complete, assuming a finite state space
- Time and space complexities are now bounded by the size of the state space
  - i.e., the number of vertices and edges,  $|V| + |E|$
- Note that we **do not** need to update under BFS and DFS since costs play no part in algorithm and they cannot guarantee an optimal solution anyway

# Questions on the Lecture?

---

- Was anything unclear?
- Do you need to clarify anything?
- Channels
  - Verbally on Zoom
  - On Archipelago
  - Via Zoom Chat



OR <https://archipelago.rocks/app/resend-invite/54568232609>