

Informed Search: Incorporating Domain Knowledge

CS3243: Introduction to Artificial Intelligence – Lecture 3

24 January 2022

Contents

1. Administrative Matters
2. Reviewing UCS
3. Greedy Best-First Search
4. A* Search
5. Dominant Heuristics

Administrative Matters

Project Poll

Grade distribution

- 20% competitive - i.e., driven by performance of the solutions submitted by your peers
 - Specifically, $\text{percentage_obtained} = 20 * [(\text{total_students} - (\text{your_rank} - 1)) / \text{total_students}]$
 - Ranks based on 1224 ranking system
- 30% tough cases - i.e., some optimisations would be required to clear
- 50% applications on standard difficulty - i.e., basic implementations

Note that only the 20% component is different between the two options.

23%

Version 1:
20% on competitive component

23%



199 responses

Grade distribution

- 20% hidden test cases - i.e., hidden cases requiring several optimisations
- 30% tough cases - i.e., some optimisations would be required to clear
- 50% applications on standard difficulty - i.e., basic implementations

Note that only the 20% component is different between the two options.

77%

Version 2:
20% on hidden test cases

77%

Project 1

- Released today
 - Individual work (no groups)
 - Python 3.7
 - Graded based on test cases (+ some inspection)
 - Public test cases given in release
 - Private test cases on *codePost*
- Deadline is 20 February 2022
 - Late penalties
 - Within deadline +24 hours = 80% of score
 - Within deadline +48 hours = 50% of score
 - Beyond deadline +48h hours = 0% of score

CNY Public Holidays

- Affected Classes
 - Monday : T02 (1300), T03 (1400), T04 (1500)
 - Tuesday: T05 (0900), T06 (1000), T07 (1100), T08 (1200)
 - Wednesday : T09 (0800), T10 (0900), T11 (1000), T12 (1100)
- Alternative Zoom Sessions
 - To be arranged by your tutors; announced next week
 - Rahul : T02, T03
 - May : T04, T07, T08
 - Bryan : T05, T06
 - Sagar : T09, T10
 - Jia Wei : T11, T12

Upcoming...

- Tutorials

- Begin this week!
 - Today: T02 (1300 hrs), T03 (1400 hrs), T04 (1500 hrs)
 - Tuesday – Friday: T05 – T17
- Face-to-face in SR-2
 - Show negative FET → uNivUS Green Pass


- Deadlines

- DQ3 (released today)
 - *Due this Friday (28 January), 2359 hrs*
- TA2 (released today)
 - *Due this Sunday (30 January), 2359 hrs*
 - *Refer to the tutorial assignment instructions document on LumiNUS*

Reviewing UCS

UCS Optimality

- UCS traverses paths in order of path cost



- When UCS pops a node from the frontier
 - Minimum path cost to that state
 - Since path costs always increase from initial state

Tree-search Versus Graph-search

```
frontier = {initial state} // frontier is a data structure
while frontier not empty:
    current = frontier.pop()
    if isGoal(current) return path found
    for a in actions(current):
        frontier.push(T(current, a))
return failure
```

UCS:

- Frontier = Priority Queue
- Priority of node n
 - Path cost of current path taken to n , $g(n)$

With a graph-search implementation:

- Maintain a *reached* hash table
- Add nodes corresponding to each state reached (i.e., on push)
- Only add new node to *frontier* (and *reached*) if
 - state represented by node not previously reached
 - path to state already reached is cheaper than one stored

UCS Under Tree-Search & Graph-Search

- Tree search will try ALL paths

- No paths excluded
- No issues with optimality


Completeness assumptions:

- b finite, and state space finite or has a solution
- All action costs are $> \epsilon > 0$

- What about graph-search?

- Ensure optimal path not among excluded paths
- Consider this example


- $F = \{A(0)\}; R = \{A\}$
 - pop $A(0)$, push $C(1)$ and $G(3)$
- $F = \{C(1), G(3)\}; R = \{A, C, G\}$
 - pop $C(1)$, push $G(2)$ since lower cost
- $F = \{G(2), G(3)\}; R = \{A, C, G\}$
 - pop $G(2)$, path is $A \rightarrow C \rightarrow G$



Notice that without the update to G while it was on the frontier, we would not have returned the optimal path

Going in the Right Direction?

- Uninformed search algorithms are systematic
 - Search outward from the initial state
 - All directions
- What can we do to try to move in the right direction?




General idea

Use domain knowledge about the problem environment to determine the cost required to go from a particular state to its nearest goal


Path Costs & Heuristics

- UCS
 - Frontier = Priority Queue
 - Priority for node $n = g(n)$
 - $g(n)$ quantifies the path cost from the initial state S to $n.state$ as defined by $n.path$
- General idea: use domain knowledge to estimate cost from $n.state$ to G
- Define a *heuristic* function h
 - h approximates the path cost from $n.state$ to its nearest goal G



Ideas on Deriving Heuristic Functions

- Consider a Maze Puzzle problem
 - Layout known
 - Moves $\leftarrow, \uparrow, \rightarrow, \downarrow$
 - Find path from **S** to **G**
- Example: Euclidean distance
 - $h(n)$ = Euclidean distance from n to **G**
- General requirements
 - Efficient – e.g., Euclidean distance is $O(m)$, where m = no. dimensions
 - More properties discussed later



$h(G) = 0$ requirement

General idea

Use domain knowledge about the costs (e.g., distances) between a given node and its closest goal – i.e., think about how to define the function h .

More on this in the next lecture.

Implementation with Evaluation Functions

- Keep using a priority queue for frontier
 - Use different priorities
- Define an evaluation function f
 - Priority for priority queue
 - Priority for node $n = f(n)$
 - UCS: $\text{priority} = f(n) = g(n)$
- Now consider different evaluation functions
 - *Greedy Best-First Search*: $\text{priority} = f(n) = h(n)$
 - *A* Search*: $\text{priority} = f(n) = g(n) + h(n)$

Best-First Search Algorithm

- General graph-search implementation

function BEST-FIRST-SEARCH(*problem*, *f*) **returns** a solution node or *failure*

node \leftarrow NODE(STATE=*problem*.INITIAL)

frontier \leftarrow a priority queue ordered by *f*, with *node* as an element

reached \leftarrow a lookup table, with one entry with key *problem*.INITIAL and value *node*

while not IS-EMPTY(*frontier*) **do**

node \leftarrow POP(*frontier*)

if *problem*.IS-GOAL(*node*.STATE) **then return** *node*

Late Goal Test

for each *child* **in** EXPAND(*problem*, *node*) **do**

s \leftarrow *child*.STATE

if *s* is not in *reached* **or** *child*.PATH-COST < *reached*[*s*].PATH-COST **then**

reached[*s*] \leftarrow *child*

add *child* to *frontier*

Graph-search

return *failure*

function EXPAND(*problem*, *node*) **yields** nodes

s \leftarrow *node*.STATE

for each *action* **in** *problem*.ACTIONS(*s*) **do**

s' \leftarrow *problem*.RESULT(*s*, *action*)

cost \leftarrow *node*.PATH-COST + *problem*.ACTION-COST(*s*, *action*, *s'*)


yield NODE(STATE=*s'*, PARENT=*node*, ACTION=*action*, PATH-COST=*cost*)

Utilises search problem definitions

Greedy Best-First Search


The Greedy Best-First Search Algorithm

- Implemented like UCS except
 - $f(n) = h(n)$
- General idea
 - Given all nodes along the frontier
 - Explore next reachable state that you estimate is closest to a goal



The Greedy Best-First Search Algorithm

- Example (tree-search)



Notice that even with the perfect heuristic, we may not get the optimal solution. Why?

Algorithm never exploits information on path already travelled.

Assume this h :

n	$h(n)$	$h^*(n)$
A	3	3
B	2	2
C	1	1
D	1	1
G	0	0

Trace:

ITR1 = $[A((-), 3)]$

ITR2 = $[C((A), 1), B((A), 2)]$

ITR3 = $[G((A, C), 0), B((A), 2)]$

ITR4 = DONE (A, C, G)

$h^*(n)$ = true path cost from n to nearest goal

Completeness & Optimality

- Tree-search version is incomplete
 - General idea
 - Can get stuck in a loop between nodes where h values are lowest
 - Prove with counter example - T02 Q1a
- Graph-search is complete as long as search space is finite
 - General idea
 - With no revisits, in finite state space, will visit entire space
 - Prove – T02 Q1b
- Not optimal under either tree-search or graph-search
 - As shown in example on last slide
 - Find another example – T02 Q1c

Questions on the Lecture so far?

- Was anything unclear?
- Do you need to clarify anything?
- Channels
 - Verbally on Zoom
 - On Archipelago
 - Via Zoom Chat




OR <https://archipelago.rocks/app/resend-invite/71722702648>

A* Search

The A* Search Algorithm


- Greedy Best-First Search
 - With greedy, $f(n) = h(n)$
 - Does not consider cost of path already taken
- Accounting for costs already incurred: A*
 - With A*, $f(n) = g(n) + h(n)$
 - $g(n)$: actual path cost from S to n
 - $h(n)$: estimated cheapest path cost from n to G
- A* priorities
 - Total path cost estimates from S to G
 - Gets more accurate as paths get explored



The A* Search Algorithm

- Example (tree-search)

same example as used on greedy (slide 19)



A* outputs the optimal solution,
unlike the Greedy Best-First Search

Will it always be optimal?
What about graph-search?

Assume this h :

again, same as before (slide 19)

n	$h(n)$	$h^*(n)$
A	3	3
B	2	2
C	1	1
D	1	1
G	0	0

Trace:

ITR1 = [A((-),0+3)]

ITR2 = [B((A),1+2), C((A),8+1)]

ITR3 = [D((A,B),2+1), C((A),8+1)]

ITR4 = [G((A,B,D),3+0), C((A),8+1)]

ITR5 = DONE (A,B,D,G)

$h^*(n)$ = true path cost from n to nearest goal

Completeness & Optimality


- Completeness
 - Same criteria as UCS
 - b finite, and state space finite or has a solution
 - All action costs $> \epsilon > 0$
- Optimality
 - Depends on the properties of h

Admissible Heuristics

- $h(n)$ is *admissible* if $\forall n, h(n) \leq h^*(n)$
 - $h(n)$ never overestimates the cost
 - Implications
 - Paths not ending at a goal are under-estimated
 - Evaluation function of value of a non-goal is under-estimated
 - At non-goal n , $f(n) = g(n) + h(n) \leq g(n) + h^*(n)$
 - Paths ending at a goal are exact
 - Evaluation function of value of a goal is exact
 - At goal m , $f(m) = g(m) + h(m)$, where $h(m) = 0$
- Examples:
 - Euclidean distance in the maze environment (always underestimates)
- Theorem: If $h(n)$ is admissible, then A^* using tree-search is optimal


Optimality of A* Using Tree-Search

- Proving the Theorem
 - T02 Q2a
- Consider the following
 - A* is optimal \rightarrow returns optimal path
 - Let s to n to t^* be the optimal path
 - If not optimal:
 - Must explore path to t first
 - Where t is a goal
 - Not optimal \rightarrow explore t before n



Optimality of Admissible A* Under Tree-Search

- Assume t expanded before n
 - $f(t) < f(n)$
- Assuming tree-search
 - All paths searched
 - All sub-paths* along a single path to a goal must be searched before that goal
 - For non-goal m ,
 $f(m) \leq g(m) + h^*(m)$ (since admissible)
 - If goal m^* on path from m ,
 $f(m) \leq f(m^*)$ (since before
 - Since t^* is goal on optimal path
 - $f(n) < f(t^*) < f(t)$
 - **CONTRADICTION**



* Consider a path, P , from an initial state s to a goal state t , to be $s > n_1 > n_2 > \dots > n_k > t$
Let a sub-path to P , P' be any path $s > n_1 > n_2 > \dots > n_i$, where $1 \leq i \leq k$

A* Using Graph-Search

- Difference between tree-search and graph-search
 - Under admissibility and tree search
 - All nodes leading to a goal are expanded before the goal
 - Optimal path will be found
 - Under graph-search we may skip some paths (due to no revisiting)
- Skipping only redundant paths
 - Graph-search checks and allows some revisits
 - As long as a path is cheaper, allow it onto the frontier even if must revisit
 - Still optimal since equivalent to tree-search


Limited-Graph-Search

- What if we just avoid revisits altogether without any exceptions?
 - i.e., as long as in reached, do not revisit, even if new paths are lower cost
- Limited-Graph-Search (version 1)
 - Just like graph-search, but no exceptions even on lower path costs
 - Uses **reached** hash table
 - Adds to reached on push to frontier
 - Only pushes to frontier when not in reached
 - Excludes **all** redundant paths, but may also exclude **some** non-redundant paths
- Limited-graph-search (version 2)
 - Similar to version 1
 - Except adds to **reached** on pop from frontier
 - Excludes **less** redundant paths than version 1*
 - Excludes **less*** non-redundant paths than version 1*

* Now allows revisits to states on the frontier, but not yet popped from the frontier

Limited-Graph-Search

- Consider limited-graph-search on UCS
- Recall UCS example
 - $F = \{A(0)\}; R = \{A\}$
 - pop $A(0)$, push $C(1)$ and $G(3)$
 - $F = \{C(1), G(3)\}; R = \{A, C, G\}$
 - pop $C(1)$, push $G(2)$
 - $F = \{G(2), G(3)\}; R = \{A, C, G\}$
 - pop $G(2)$, path is $A \rightarrow C \rightarrow G$



This works only under limited-graph-search version 2, and not version 1, for a similar reason to why an Early Goal Test would cause UCS to not return an optimal solution

From this point, let limited-graph-search imply limited-graph-search version 2. We will not study version 1 any further


Limited-Graph-Search Consistent Heuristics

- Does this mean A^* is optimal under limited-graph-search?
- Example

T02 Q3b - construct an alternative example

Assume this admissible h :

n	$h(n)$	$h^*(n)$
S	8	8
A	7	7
B	0	6
G	0	0



Trace:

ITR1 = $[S((-), 0+8)]$

ITR2 = $[B((S), 3+0), A((S), 1+7)]$

ITR3 = $[A((S), 1+7), G((S,B), 9+0)]$

ITR4 = $[G((S,B), 9+0)]$ as B popped before, do not revisit

ITR5 = DONE (S,B,G) not the optimal path!

We need a tighter constraint on h


Similar to what UCS offers → contours of search progression

Why Not Optimal?

- Consider the previous example

Assume this
admissible h :

n	$h(n)$	$h^*(n)$
S	8	8
A	7	7
B	0	6
G	0	0



Observe the sequence of $f(n) = g(n) + h(n)$ values along each path:

path to n (from S)	$g(n) + h(n)$	$g(n) + h^*(n)$
S	$0+8$	$0+8$
$S > A$	$1+7$	$1+7$
$S > A > B$	$2+0$	$2+6$
$S > A > B > G$	$8+0$	$8+0$

Dip!

path to n (from S)	$g(n) + h(n)$	$g(n) + h^*(n)$
S	$0+8$	$0+9$
$S > B$	$3+0$	$3+6$
$S > B > G$	$9+0$	6

Dip!

Consistent Heuristics

- Forming contours
 - Under tree-search g costs are monotonically increasing
 - For f costs to be monotonically increasing along a path
 - We need: $g(n) + h(n) \leq g(n) + \text{cost}(n, a, n') + h(n')$
 - And thus $h(n) \leq \text{cost}(n, a, n') + h(n')$
 - We will use the above requirement

- $h(n)$ is **consistent** if $\forall n$, and successor of n, n' ,
$$h(n) \leq \text{cost}(n, a, n') + h(n')$$

Note that:
consistency \Rightarrow admissibility
Proof – T02 Q3a

- Theorem: If $h(n)$ is consistent, then A^* using graph-search is optimal
- Prove this in a similar manner to the UCS proof (contours) – T02 Q2b

Dominant Heuristics

Efficiency & Dominance

- Efficiency of A^* depends on the accuracy of its heuristics
 - Higher heuristic accuracy means we need to try fewer paths
 - Specifics not covered in CS3243
- Which heuristics are better?
- If $h_1(n) \geq h_2(n)$ for all n , then h_1 ***dominates*** h_2
 - If h_1 is also *admissible*
 - h_1 must be closer to h^* than h_2
 - h_1 must be more efficient than h_2

Note: with some interpretations, dominance requires admissibility. We apply a more generic version that does not.

Questions on the Lecture?

- Was anything unclear?
- Do you need to clarify anything?
- Channels
 - Verbally on Zoom
 - On Archipelago
 - Via Zoom Chat



OR <https://archipelago.rocks/app/resend-invite/75289652625>